

ECOCIVIX AI

09 — Safety & Trust Specification

TEAM VELOCITY • HackYatra × EDM 3.0 • PRE-HACKATHON PLANNING

Purpose: Define responsible boundaries for safety, privacy, AI uncertainty and user trust so ECOCIVIX does not make unsupported claims.

Status: PRE-HACKATHON DRAFT / PLANNING SPECIFICATION

Approval authority: Garv Nain

Safety Position

- Safety is a focused module within civic intelligence, not a separate surveillance or crime-prediction platform.

Allowed Safety Use

- Safety-related civic reporting; safety relevance of infrastructure conditions; emergency-support workflow where actually implemented; emergency contacts; optional location sharing with chosen contacts where supported.

Prohibited Claims

- No guaranteed protection; no guaranteed emergency response; no crime prediction; no perfect threat detection; no unsupported claims about law-enforcement outcomes.

AI Limitations

- AI assessments are probabilistic recommendations and can be wrong. The UI must not represent AI as an authority that guarantees outcomes.

Human Oversight

- Administrators remain responsible for operational decisions. AI should assist prioritization and routing, not silently make irreversible decisions.

Location Privacy

- Explain location use when requesting it. Store only what is needed. Avoid exposing precise personal locations publicly unless the product and user consent explicitly support it.

Data Minimization

- Collect only information needed for reporting, resolution, authentication and legitimate community intelligence.

Access Control

- Citizens should only access resources they are authorized to view. Admin operations require server-side role checks.

File Safety

- Validate uploaded file type/size where images are supported. Store through controlled storage and avoid exposing private files unintentionally.

Auditability

- Important admin actions should be represented in the issue timeline/activity system where practical.

Transparency

- Clearly label AI-generated assessment, human administrative decision, and citizen-submitted information.

Emergency Handling

- If emergency support is included, direct users to appropriate real-world emergency services/contacts rather than claiming ECOCIVIX itself provides emergency response.

Acceptance Criteria

- No unsupported safety claims; privacy-aware location handling; role-based authorization; transparent AI labeling; secure file handling; meaningful audit trail.